

ClickHouse Cheat Sheet

Quick Start Commands

Connection

```
# Connect via clickhouse-client
clickhouse-client --host localhost --port 9000 --user default

# Connect via HTTP
curl 'http://localhost:8123/' --data-binary "SELECT version()"

# Connect with password
clickhouse-client --host localhost --port 9000 --user myuser --password mypass
```

Basic Database Operations

```
-- Show databases
SHOW DATABASES;

-- Create database
CREATE DATABASE my_db;

-- Use database
USE my_db;

-- Drop database
DROP DATABASE my_db;
```

Essential Configuration

Docker Compose Setup

```
version: '3.8'
services:
  clickhouse:
    image: clickhouse/clickhouse-server:latest
    container_name: clickhouse-server
    ports:
      - "8123:8123" # HTTP
      - "9000:9000" # Native TCP
    volumes:
      - ./data:/var/lib/clickhouse
      - ./logs:/var/log/clickhouse-server
```

```

- ./config:/etc/clickhouse-server/config.d
environment:
- CLICKHOUSE_USER=default
- CLICKHOUSE_PASSWORD=
- CLICKHOUSE_DEFAULT_ACCESS_MANAGEMENT=1
ulimits:
  nofile:
    soft: 262144
    hard: 262144

```

Critical Server Settings (`config.xml`)

```

<clickhouse>
  <!-- Memory -->

  <max_server_memory_usage_to_ram_ratio>0.9</max_server_memory_usage_to_ram_ratio>
  <max_memory_usage>10000000000</max_memory_usage>

  <!-- Connections -->
  <max_concurrent_queries>100</max_concurrent_queries>
  <max_connections>4096</max_connections>

  <!-- Background processing -->
  <background_pool_size>16</background_pool_size>

  <background_merges_mutations_concurrency_ratio>2</background_merges_mutations_concurrency_ratio>

  <!-- Compression -->
  <compression>
    <case><method>lz4</method></case>
  </compression>
</clickhouse>

```

Performance Settings

```

-- Memory management
SET max_memory_usage = 10000000000;           -- 10GB per query
SET max_bytes_before_external_group_by = 20000000000; -- 20GB before
spilling
SET max_bytes_before_external_sort = 20000000000;   -- 20GB before
spilling

-- Query optimization
SET max_threads = 8;                               -- Parallel threads
SET max_execution_time = 3600;                      -- 1 hour timeout
SET optimize_read_in_order = 1;                    -- Read
optimization

```

```

SET optimize_aggregation_in_order = 1;           -- GROUP BY
optimization

-- Insert optimization
SET max_insert_block_size = 1048576;             -- 1M rows per
block
SET async_insert = 1;                             -- Enable async
inserts
SET async_insert_timeout_ms = 200;                -- 200ms timeout

-- JOIN optimization
SET join_algorithm = 'hash';                      -- Hash join
(default)
SET max_bytes_in_join = 1000000000;              -- 1GB join limit
SET join_use_nulls = 1;                          -- Use NULL for
non-matches

```

Data Types Quick Reference

Numeric Types

```

-- Integers
Int8, Int16, Int32, Int64, Int128, Int256        -- Signed integers
UInt8, UInt16, UInt32, UInt64, UInt128, UInt256  -- Unsigned
integers

-- Floating point
Float32, Float64                                -- IEEE 754
floating point

-- Fixed precision
Decimal32(scale), Decimal64(scale), Decimal128(scale) -- Fixed decimal
precision

```

String & Text Types

```

String                                           -- Variable-length
UTF-8
FixedString(N)                                  -- Fixed-length (N
bytes)
LowCardinality(String)                          -- Dictionary-
encoded strings
UUID                                             -- 128-bit UUID

```

Date & Time Types

Date	-- YYYY-MM-DD
Date32	-- Extended range
date	
DateTime	-- Unix timestamp
DateTime('timezone')	-- Timezone-aware
timestamp	
DateTime64(precision)	-- High-precision
timestamp	

Complex Types

Array(T)	-- Array of type T
Tuple(T1, T2, ...)	-- Fixed-size tuple
Map(K, V)	-- Key-value map
Nested(col1 T1, col2 T2, ...)	-- Nested structure
Enum8('val1'=1, 'val2'=2)	-- 8-bit enum
Nullable(T)	-- Nullable type (use
sparingly)	

Table Engines

MergeTree Family

```
-- Basic MergeTree
CREATE TABLE events (
    id UInt64,
    timestamp DateTime,
    user_id UInt32,
    event_type String
) ENGINE = MergeTree()
PARTITION BY toYYYYMM(timestamp)
ORDER BY (user_id, timestamp);

-- ReplacingMergeTree (deduplication)
CREATE TABLE user_profiles (
    user_id UInt32,
    name String,
    updated_at DateTime
) ENGINE = ReplacingMergeTree(updated_at)
ORDER BY user_id;

-- CollapsingMergeTree (updates/deletes)
CREATE TABLE user_actions (
    user_id UInt32,
    action String,
    timestamp DateTime,
    sign Int8
) ENGINE = CollapsingMergeTree(sign)
```

```
ORDER BY (user_id, timestamp);

-- SummingMergeTree (pre-aggregation)
CREATE TABLE daily_stats (
    date Date,
    user_id UInt32,
    page_views UInt64,
    clicks UInt64
) ENGINE = SummingMergeTree((page_views, clicks))
PARTITION BY toYYYYMM(date)
ORDER BY (date, user_id);
```

Distributed Tables

```
-- Create distributed table
CREATE TABLE events_distributed AS events
ENGINE = Distributed(cluster_name, database_name, table_name,
sharding_key);
```

Essential SQL Operations

CREATE TABLE Patterns

```
-- Analytics table with partitioning
CREATE TABLE analytics_events (
    event_id UUID,
    timestamp DateTime,
    user_id UInt32,
    session_id String,
    event_type LowCardinality(String),
    properties Map(String, String),

    -- Indexes
    INDEX event_type_idx event_type TYPE bloom_filter GRANULARITY 1,
    INDEX user_id_idx user_id TYPE minmax GRANULARITY 1
) ENGINE = MergeTree()
PARTITION BY toYYYYMM(timestamp)
ORDER BY (user_id, timestamp, event_id)
SETTINGS
    index_granularity = 8192,
    storage_policy = 'default';
```

INSERT Patterns

```
-- Single insert
INSERT INTO events VALUES (1, now(), 1001, 'click');
```

```

-- Batch insert
INSERT INTO events VALUES
    (1, now(), 1001, 'click'),
    (2, now(), 1002, 'view'),
    (3, now(), 1003, 'purchase');

-- Insert from SELECT
INSERT INTO events_backup
SELECT * FROM events
WHERE timestamp >= '2024-01-01';

-- Insert with functions
INSERT INTO events
SELECT
    generateUUIDv4(),
    now() - toIntervalDay(number % 30),
    toUInt32(rand() % 10000),
    'generated_event'
FROM numbers(1000000);

```

SELECT Patterns

```

-- Basic filtering
SELECT * FROM events
WHERE event_type = 'click'
    AND timestamp >= '2024-01-01'
LIMIT 100;

-- Aggregation
SELECT
    event_type,
    count() as event_count,
    uniq(user_id) as unique_users,
    avg(session_duration) as avg_duration
FROM events
GROUP BY event_type
ORDER BY event_count DESC;

-- Time-based aggregation
SELECT
    toStartOfDay(timestamp) as day,
    count() as daily_events
FROM events
WHERE timestamp >= now() - INTERVAL 7 DAY
GROUP BY day
ORDER BY day;

-- Window functions
SELECT
    user_id,

```

```

        timestamp,
        event_type,
        row_number() OVER (PARTITION BY user_id ORDER BY timestamp) as
event_sequence
FROM events
ORDER BY user_id, timestamp;

```

JOIN Patterns

```

-- INNER JOIN
SELECT
    e.event_type,
    u.username,
    e.timestamp
FROM events e
INNER JOIN users u ON e.user_id = u.user_id
WHERE e.timestamp >= '2024-01-01';

-- LEFT JOIN with aggregation
SELECT
    u.username,
    count(e.event_id) as event_count
FROM users u
LEFT JOIN events e ON u.user_id = e.user_id
GROUP BY u.username
ORDER BY event_count DESC;

```

Indexing & Optimization

Secondary Indexes

```

-- Bloom filter (for equality checks)
ALTER TABLE events ADD INDEX event_type_bloom event_type TYPE bloom_filter
GRANULARITY 1;

-- Set index (for IN queries)
ALTER TABLE events ADD INDEX status_set status TYPE set(0) GRANULARITY 1;

-- MinMax index (for range queries)
ALTER TABLE events ADD INDEX timestamp_minmax timestamp TYPE minmax
GRANULARITY 1;

-- N-gram index (for text search)
ALTER TABLE events ADD INDEX content_ngram content TYPE ngrambf_v1(3, 256,
2, 0) GRANULARITY 1;

```

Materialized Views

```

-- Create materialized view for aggregation
CREATE MATERIALIZED VIEW daily_user_stats
ENGINE = SummingMergeTree()
PARTITION BY toYYYYMM(date)
ORDER BY (date, user_id)
AS SELECT
    toDate(timestamp) as date,
    user_id,
    count() as event_count,
    uniq(session_id) as session_count
FROM events
GROUP BY date, user_id;

```

Projections

```

-- Add projection for different sort order
ALTER TABLE events
ADD PROJECTION timestamp_projection (
    SELECT * ORDER BY timestamp, user_id
);

-- Materialize projection
ALTER TABLE events MATERIALIZE PROJECTION timestamp_projection;

```

Essential Functions

Date/Time Functions

now()	-- Current timestamp
today()	-- Current date
yesterday()	-- Yesterday's date
toYYYYMM(date)	-- Extract YYYYMM
toStartOfMonth(date)	-- First day of month
toStartOfWeek(date)	-- Start of week
toStartOfDay(date)	-- Start of day
toStartOfHour(date)	-- Start of hour
dateDiff('day', date1, date2)	-- Date difference
dateAdd('day', 7, date)	-- Add days to date
formatDateTime(date, '%Y-%m-%d %H:%M:%S')	-- Format date

String Functions

length(string)	-- String length
lower(string), upper(string)	-- Case conversion
substring(string, start, length)	-- Extract substring

<code>concat(str1, str2, ...)</code>	-- Concatenate strings
<code>splitByChar(',', string)</code>	-- Split string
<code>replaceAll(string, 'old', 'new')</code>	-- Replace all occurrences
<code>trim(string)</code>	-- Remove whitespace

Aggregate Functions

<code>count()</code>	-- Count rows
<code>count(DISTINCT column)</code>	-- Count unique values
<code>sum(column)</code>	-- Sum values
<code>avg(column)</code>	-- Average
<code>min(column), max(column)</code>	-- Min/max values
<code>uniq(column)</code>	-- Approximate unique count
<code>uniqExact(column)</code>	-- Exact unique count
<code>quantile(0.5)(column)</code>	-- Median (50th percentile)
<code>quantiles(0.25, 0.5, 0.75)(column)</code>	-- Multiple quantiles
<code>groupBy(column)</code>	-- Collect values into array

Array Functions

<code>arrayJoin(array)</code>	-- Expand array to rows
<code>arrayMap(x -> x * 2, array)</code>	-- Apply function to elements
<code>arrayFilter(x -> x > 0, array)</code>	-- Filter array elements
<code>arrayReduce('sum', array)</code>	-- Reduce array
<code>arraySort(array)</code>	-- Sort array
<code>arrayReverse(array)</code>	-- Reverse array
<code>has(array, element)</code>	-- Check if element exists
<code>length(array)</code>	-- Array length

Type Conversion Functions

<code>toString(value)</code>	-- Convert to string
<code>toInt32(value)</code>	-- Convert to integer
<code>toFloat64(value)</code>	-- Convert to float
<code>toDate(string)</code>	-- Convert to date
<code>toDateTime(string)</code>	-- Convert to datetime
<code>toUUID(string)</code>	-- Convert to UUID
<code>CAST(value AS Type)</code>	-- Generic type casting

Monitoring & Troubleshooting

System Tables

```

-- Current processes
SELECT query_id, user, query, elapsed, memory_usage
FROM system.processes
WHERE query != '';

-- Query log (recent queries)
SELECT
    query_start_time,
    query_duration_ms,
    query,
    read_rows,
    memory_usage
FROM system.query_log
WHERE event_time > now() - INTERVAL 1 HOUR
    AND type = 'QueryFinish'
ORDER BY query_duration_ms DESC
LIMIT 10;

-- Table sizes
SELECT
    database,
    table,
    sum(rows) as total_rows,
    formatReadableSize(sum(bytes_on_disk)) as size
FROM system.parts
WHERE active = 1
GROUP BY database, table
ORDER BY sum(bytes_on_disk) DESC;

-- Parts information
SELECT
    table,
    partition,
    name,
    rows,
    formatReadableSize(bytes_on_disk) as size,
    modification_time
FROM system.parts
WHERE active = 1 AND table = 'my_table'
ORDER BY modification_time DESC;

```

Performance Monitoring

```

-- Slow queries
SELECT
    query,
    query_duration_ms,
    read_rows,
    read_bytes,
    memory_usage

```

```

FROM system.query_log
WHERE query_duration_ms > 5000
    AND event_time > now() - INTERVAL 1 DAY
ORDER BY query_duration_ms DESC
LIMIT 20;

-- Memory usage
SELECT
    event_time,
    CurrentMetric_MemoryTracking as memory_usage,
    CurrentMetric_BackgroundPoolTask as background_tasks
FROM system.metric_log
WHERE event_time > now() - INTERVAL 1 HOUR
ORDER BY event_time DESC
LIMIT 100;

-- Merge performance
SELECT
    table,
    count() as merge_count,
    avg(duration_ms) as avg_duration,
    sum(bytes_read_uncompressed) as total_bytes
FROM system.part_log
WHERE event_type = 'MergeParts'
    AND event_time > now() - INTERVAL 1 DAY
GROUP BY table
ORDER BY avg_duration DESC;

```

Health Checks

```

-- Cluster status
SELECT
    'Version' as metric, version() as value
UNION ALL
SELECT 'Uptime', toString(uptime())
UNION ALL
SELECT 'Active Parts', toString(count()) FROM system.parts WHERE active =
1;

-- Replication status (for replicated tables)
SELECT
    database,
    table,
    replica_name,
    log_max_index - log_pointer as replication_lag
FROM system.replicas;

-- Disk usage
SELECT
    name,
    formatReadableSize(free_space) as free_space,

```

```
formatReadableSize(total_space) as total_space,  
round((free_space / total_space) * 100, 2) as free_percent  
FROM system.disks;
```

Common Patterns & Best Practices

Partitioning Strategies

```
-- Monthly partitioning (most common)  
PARTITION BY toYYYYMM(timestamp)  
  
-- Daily partitioning (high volume)  
PARTITION BY toYYYYMMDD(timestamp)  
  
-- Custom partitioning  
PARTITION BY (toYYYYMM(timestamp), user_type)
```

ORDER BY Best Practices

```
-- High to low cardinality  
ORDER BY (user_id, timestamp, event_id) -- Good  
ORDER BY (timestamp, user_id, event_id) -- Less optimal for user queries  
  
-- Include frequently filtered columns  
ORDER BY (status, created_date, id) -- If you often filter by status
```

Compression Settings

```
-- Table with compression codecs  
CREATE TABLE compressed_logs (  
    timestamp DateTime CODEC(Delta, LZ4),  
    user_id UInt32 CODEC(LZ4),  
    message String CODEC(ZSTD(3)),  
    level Enum8('DEBUG'=1, 'INFO'=2, 'ERROR'=3) CODEC(LZ4)  
) ENGINE = MergeTree()  
ORDER BY (timestamp, user_id);
```

Batch Processing

```
-- Optimal settings for batch inserts  
SET max_insert_block_size = 1048576;  
SET async_insert = 1;  
SET async_insert_timeout_ms = 200;
```

```
-- Generate test data
INSERT INTO events
SELECT
    generateUUIDv4(),
    now() - toIntervalSecond(number),
    toUInt32(rand() % 10000),
    ['click', 'view', 'purchase'][rand() % 3 + 1]
FROM numbers(1000000);
```

Troubleshooting Commands

Query Optimization

```
-- Explain query plan
EXPLAIN SELECT * FROM events WHERE user_id = 1001;

-- Explain with pipeline
EXPLAIN PIPELINE SELECT count() FROM events GROUP BY user_id;

-- Check if indexes are used
SELECT * FROM system.query_log
WHERE query LIKE '%your_query%'
AND ProfileEvents['SelectedMarks'] > 0;
```

Memory Issues

```
-- Check memory usage by query
SELECT
    query_id,
    memory_usage,
    peak_memory_usage,
    query
FROM system.query_log
WHERE memory_usage > 1000000000 -- > 1GB
ORDER BY memory_usage DESC;

-- Current memory usage
SELECT formatReadableSize(sum(memory_usage)) as total_memory
FROM system.processes;
```

Performance Tuning

```
-- Settings for different workloads

-- OLAP/Analytics workload
SET max_memory_usage = 20000000000;
```

```

SET max_threads = 16;
SET optimize_read_in_order = 1;
SET compile_expressions = 1;

-- High-throughput inserts
SET max_insert_block_size = 1048576;
SET async_insert = 1;
SET max_insert_threads = 8;

-- Memory-constrained environment
SET max_memory_usage = 2000000000;
SET max_bytes_before_external_group_by = 1000000000;
SET join_algorithm = 'partial_merge';

```

Quick Reference Commands

Administrative

SHOW TABLES;	-- List tables
DESCRIBE table_name;	-- Table schema
SHOW CREATE TABLE table_name;	-- Table DDL
OPTIMIZE TABLE table_name FINAL;	-- Force merge
TRUNCATE TABLE table_name;	-- Delete all data
DROP TABLE table_name;	-- Delete table

Useful Shortcuts

```

-- Current settings
SELECT name, value FROM system.settings WHERE changed = 1;

-- Table statistics
SELECT count(), min(timestamp), max(timestamp) FROM events;

-- Sample data
SELECT * FROM events SAMPLE 0.1 LIMIT 10; -- 10% sample

-- Check data types
SELECT name, type FROM system.columns WHERE table = 'events';

```

This cheat sheet covers the most commonly used ClickHouse operations and configurations. Keep it handy for quick reference during development and operations.